

Flight and Hotel Price Analysis from Kayak.ae

توفيق كامل توفيق حسين 23011240

ساهر سامي 23011280

نور الدين امجد عبد الفتاح احمد محمد 23011596

Brief Description:

This project involves scraping data from Kayak.ae, a travel search engine that aggregates prices for flights, hotels, and car rentals. The focus is on extracting flight and hotel pricing(**BLOCKED**) data for various destinations and analyzing how prices vary based on multiple factors.

Goal of the Project:

- Analyze pricing trends for flights and hotels based on travel dates, destinations, and booking time.
- Answer questions such as:
 - How do prices vary by day of the week or booking time?
 - What are the most affordable destinations from a specific city?
 - Are there noticeable seasonal trends?

Data Extraction Approach:

- Use Python with Selenium to scrape dynamic content.

```

1 from selenium.webdriver.common.by import By # for locating Elements
2 from selenium import webdriver #for controlling web browser
3 from selenium.webdriver.chrome.service import Service as ChromeService # using chrome as web browser
4 from webdriver_manager.chrome import ChromeDriverManager # for installing chromeDriver
5 from selenium.webdriver.support.ui import WebDriverWait #used to wait for web content to showup
6 from selenium.common.exceptions import NoSuchElementException, TimeoutException #error handling
7 import time

```

We use chrome driver to simulate web browser to scrape the flight data

- Simulate user searches for routes and hotel stays(**BLOCKED**).

```

1 departure_date = "2025-04-23"
2 return_flight_date = "2025-04-30"
3
4 routes = [
5     ["CAI", "SVO", departure_date, return_flight_date], # Cairo → Moscow
6     ["CAI", "LAX", departure_date, return_flight_date], # Cairo → Los Angeles
7     ["CAI", "MNL", departure_date, return_flight_date], # Cairo → Manila
8     ["CAI", "SGN", departure_date, return_flight_date], # Cairo → Ho Chi Minh
9     ["CAI", "SZX", departure_date, return_flight_date], # Cairo → Shenzhen
10    ["CAI", "TPE", departure_date, return_flight_date], # Cairo → Taipei
11    ["CAI", "YVR", departure_date, return_flight_date], # Cairo → Vancouver
12    ["CAI", "LIM", departure_date, return_flight_date], # Cairo → Lima
13    ["CAI", "BOG", departure_date, return_flight_date], # Cairo → Bogotá

```

Sample of routes that we used in the scrapping in the full routes file it contain 102 routes

- Extract destination, price, date, airline/hotel name, and rating(**BLOCKED**)

```

1 def scrape_kayak_flights(origin, destination, depart_date, return_date):
2     base_url = "https://www.kayak.ae/flights"
3     search_url = f"{base_url}/{origin}-{destination}/{depart_date}/{return_date}?sort=price_a"
4     print(f"Navigating to: {search_url}")
5
6     options = webdriver.ChromeOptions()
7     # options.add_argument('--headless')
8     options.add_argument('--no-sandbox')
9     options.add_argument('--disable-dev-shm-usage')
10    options.add_argument("user-agent=Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/108.0.0.0 Safari/537.36")
11
12    driver = None
13    scraped_data = []

```

We start with defining the scrape function which take the origin, destination, departure date and return date

In the second part we define the url that will be scraped and options for the chrome driver

```

1 try:
2     print("Initializing WebDriver...")
3     driver = webdriver.Chrome(service=ChromeService(ChromeDriverManager().install()), options=options)
4     driver.implicitly_wait(5)
5     print("WebDriver initialized.")
6     driver.get(search_url)
7
8     # Wait for results
9     individual_results_selector_css = "div[class*='Fwx9-result-item-container']"
10    print(f"Waiting for flight results using CSS selector: '{individual_results_selector_css}'...")
11    wait = WebDriverWait(driver, 5)
12    wait.until(EC.presence_of_element_located((By.CSS_SELECTOR, individual_results_selector_css)))
13    print("At least one result item appears to be present. Waiting for dynamic loading...")
14
15    # Delay for dynamic content
16    time.sleep(5)
17
18    # Find result elements
19    result_elements = driver.find_elements(By.CSS_SELECTOR, individual_results_selector_css)
20    print(f"Found {len(result_elements)} potential flight results.")
21
22    if not result_elements:
23        print("No flight result elements found. Check the selector or page structure.")
24        return []

```

In this code, we initialize a Chrome WebDriver using ChromeDriverManager to automatically handle the driver installation. After setting an implicit wait, we navigate to the target URL (search_url).

We then define a CSS selector to identify the container holding

individual flight results. To ensure the page is fully loaded, we explicitly wait until at least one matching element appears in the DOM.

Afterward, we add an extra short sleep to allow any dynamic content to finish rendering.

Finally, we search for all elements matching the flight result container selector. If no results are found, a message is printed to indicate that the selector might be incorrect or the page structure may have changed.

```
1 count = 0
2     for result in result_elements:
3         count += 1
4         print(f"\n--- Processing Result {count} ---")
5         try:
6             # Price
7             price_selector_css = "div[class*='e2GB-price-text']"
8             try:
9                 price = result.find_element(By.CSS_SELECTOR, price_selector_css).text.strip()
10                print(f"    Price: {price}")
11            except NoSuchElementException:
12                print(f"    - Warning: Price not found with selector: {price_selector_css}")
13                price = "N/A"
14
15            # Airlines
16            airline_selector_css = "div[class*='c5iUd-leg-carrier'] img"
17            try:
18                airline_element = result.find_element(By.CSS_SELECTOR, airline_selector_css)
19                airline = airline_element.get_attribute("alt")
20                print(f"    Airlines: {airline}")
21            except NoSuchElementException:
22                print(f"    - Warning: Airlines not found with selector: {airline_selector_css}")
23                airline = "N/A"
```

After retrieving the container that holds all the flight data, we iterate through each result element.

For each flight entry, we first attempt to extract the price by locating the appropriate sub-element using its CSS selector. If the price is successfully found, we extract and clean its text; if not, we assign "N/A" as a fallback value.

Similarly, we attempt to extract the airline name by finding the airline logo element and reading its alt attribute. If the airline information cannot be found, we also assign "N/A" as a default.

This ensures that missing fields do not cause the script to crash and that each flight entry remains consistent even if some data is unavailable.

```
1 stops_selector_css = "span[class*='JWE0-stops-text']"
2 try:
3     stops_text = result.find_element(By.CSS_SELECTOR, stops_selector_css)
4     stops = stops_text.text.strip()
5     print(f" Stops: {stops}")
6 except NoSuchElementException:
7     print(f" - Warning: Stops not found with selector: {stops_selector_css}")
8     stops = "N/A"
9
10 # Duration
11 duration_selector_css = "div[class*='xdw8 xdw8-mod-full-airport']"
12 try:
13     duration = result.find_element(By.CSS_SELECTOR, duration_selector_css).text.strip()
14     print(f" Duration: {duration}")
15 except NoSuchElementException:
16     print(f" - Warning: Duration not found with selector: {duration_selector_css}")
17     duration = "N/A"
18
19 flight_info = {
20     "origin": origin,
21     "destination": destination,
22     "price": price,
23     "airlines": airline,
24     "stops": stops,
25     "duration": duration,
26     "search_url": search_url
27 }
28 scraped_data.append(flight_info)
29 print(f" => Added: {airline} - {price} - {stops} - {duration}")
30
31 except NoSuchElementException as e:
32     print(f" - Warning: Issue processing result {count}. Skipping. Error: {e}")
33     continue
34 except Exception as e:
35     print(f" - Error processing result {count}: {e}")
36     continue
37
```

continuing as above but we include flights stops we attempt to extract the stops by locating the appropriate sub-element using its CSS selector. If the stops is successfully found, we extract and clean its text; if not, we assign "N/A" as a fallback value
And if there is any complication during the process it will throw an exception
And finally, we store all our finding into a dictionary so we can process it later

```

flights = [] # extracting data
print("--- Starting Kayak Flight Scraper ---")
for route in routes:
    flight = scrape_kayak_flights(*route)
    flights.append(flight)

```

```

flattened = [item for sublist in flights for item in sublist] # preparing to save in csv
df = pd.DataFrame(flattened)

```

```
df.head()
```

	origin	destination	price	airlines	stops	duration	search_url
0	JFK	BUD	AED 2,498	United Airlines	1 stop	13h 10m	https://www.kayak.ae/flights/JFK-BUD/2025-04-2...
1	JFK	BUD	AED 2,519	SWISS	1 stop	10h 30m	https://www.kayak.ae/flights/JFK-BUD/2025-04-2...
2	JFK	BUD	AED 2,519	SWISS	1 stop	11h 35m	https://www.kayak.ae/flights/JFK-BUD/2025-04-2...
3	JFK	BUD	AED 2,527	United Airlines	1 stop	13h 10m	https://www.kayak.ae/flights/JFK-BUD/2025-04-2...
4	JFK	BUD	N/A	N/A	N/A	N/A	https://www.kayak.ae/flights/JFK-BUD/2025-04-2...

```
df.to_csv("flights.csv", index = False)
```

combining different runs of data extraction

```

df_flights_1 = pd.read_csv("flights.csv")
df_flights_2 = pd.read_csv("flights_1.csv")
df_flights_3 = pd.read_csv("flights_2.csv")

```

```
df_flights = pd.concat([df_flights_1, df_flights_2, df_flights_3])
```

```
df_flights.to_csv("AllFlights.csv", index = False)
```

Running the script multiple time and saving the all results in AllFlights.csv

Data Cleaning:

- Normalize dates and prices.
- Handle missing or inconsistent entries.

```
# dropping search_url
df.drop(columns = 'search_url', inplace = True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1214 entries, 0 to 1213
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   origin          1214 non-null   object
 1   destination      1214 non-null   object
 2   price            735 non-null    object
 3   airlines         735 non-null    object
 4   stops            735 non-null    object
 5   duration         735 non-null    object
dtypes: object(6)
memory usage: 57.0+ KB
```

```
df.drop_duplicates(inplace = True)
```

After reading the csv file into pandas data frame we are dropping the Search_url because it unnecessary column and dropping duplicates

```
df.isna().sum() #cheaking for NAN values
```

```
origin          0
destination     0
price           80
airlines        80
stops           80
duration        80
dtype: int64
```

```
df.isnull().sum() # cheaking for nulls
```

```
origin          0
destination     0
price           80
airlines        80
stops           80
duration        80
dtype: int64
```

```
df.dropna(inplace = True)
```

```
df.shape
```

```
(280, 6)
```

Checking for Null or NaN values and dropping them

```
#converting price to USD
def convert_to_usd(price):
    only_price = price[4:].replace(',', '')
    price_in_usd = int(only_price) * 0.27 #converting the currency
    return price_in_usd

df['price'] = df['price'].apply(convert_to_usd)
```

Using the latest exchange rate from United Arab Emirates Dirham to USD,
we convert the price

- Convert text to categorical/numerical formats.


```

# converting duration to minutes form
# extracting hour and minutes using regex

def convert_to_minutes(duration):
    match = re.search(r'(\d+)h\s+(\d+)m', duration) # separating hours and minutes
    if match:
        hours = int(match.group(1)) if match.group(1) else 0
        minutes = int(match.group(2)) if match.group(2) else 0
        duration_in_minutes = hours * 60 + minutes # covering to minutes
        return duration_in_minutes
    return None

df['duration'] = df['duration'].apply(convert_to_minutes)

```

Using regex, we successfully extracted hours and minutes from the duration string and converted them into total minutes

```

#converting stops to int
def convert_to_int(stops):
    return 0 if stops == 'direct' else int(stops[0]) # if it direct it will return zero

df['stops'] = df['stops'].apply(convert_to_int)

```

Covertng stops to int and if it direct replace it with zero

```
[15]: df.head()
```

```
[15]:
```

	origin	destination	price	airlines	stops	duration
0	HBE	MLA	1058.40	Turkish Airlines	1	425
1	HBE	MLA	1058.40	Turkish Airlines	1	930
4	HBE	PRG	519.48	طيران الجزيرة	2	1335
5	HBE	PRG	545.67	طيران الجزيرة	2	1335
6	HBE	PRG	611.82	Air Arabia	2	1375

```

[ ]: #saving the result
df.to_csv("data.csv", index = False)

```

The data after it cleaned

Analysis & Visualization:

- Use Pandas for EDA (Exploratory Data Analysis).

```
: df.describe()
```

```
:
```

	price	stops	duration
count	280.000000	280.000000	280.000000
mean	1096.898143	1.196429	1207.900000
std	559.552611	0.656362	698.913646
min	129.060000	0.000000	75.000000
25%	687.757500	1.000000	672.500000
50%	1066.365000	1.000000	1105.000000
75%	1431.675000	2.000000	1601.250000
max	2952.450000	3.000000	3105.000000

```
[55]: len(df["destination"].unique())
```

```
[55]: 70
```

```
[56]: len(df['airlines'].unique())
```

```
[56]: 33
```

Price:

- Mean: 1,096.90
- Range: 129.06 – 2,952.45

Stops:

- Most flights have 1 or 2 stops

Destinations:

- 70 unique destinations


Airlines:

- 33 unique airlines

- Visualize with Matplotlib and Seaborn.

```
top_destinations = df["destination"].value_counts().head(10).index
top_avg_price_per_dest = df[df["destination"].isin(top_destinations)].groupby("destination")["price"].mean().sort_values(ascending = False)

sns.barplot(x = top_avg_price_per_dest.index,
            hue = top_avg_price_per_dest.index,
            y = top_avg_price_per_dest.values, palette= "tab10", legend = False)
plt.title("Top 10 Average Price per Destinations")
plt.show()
```



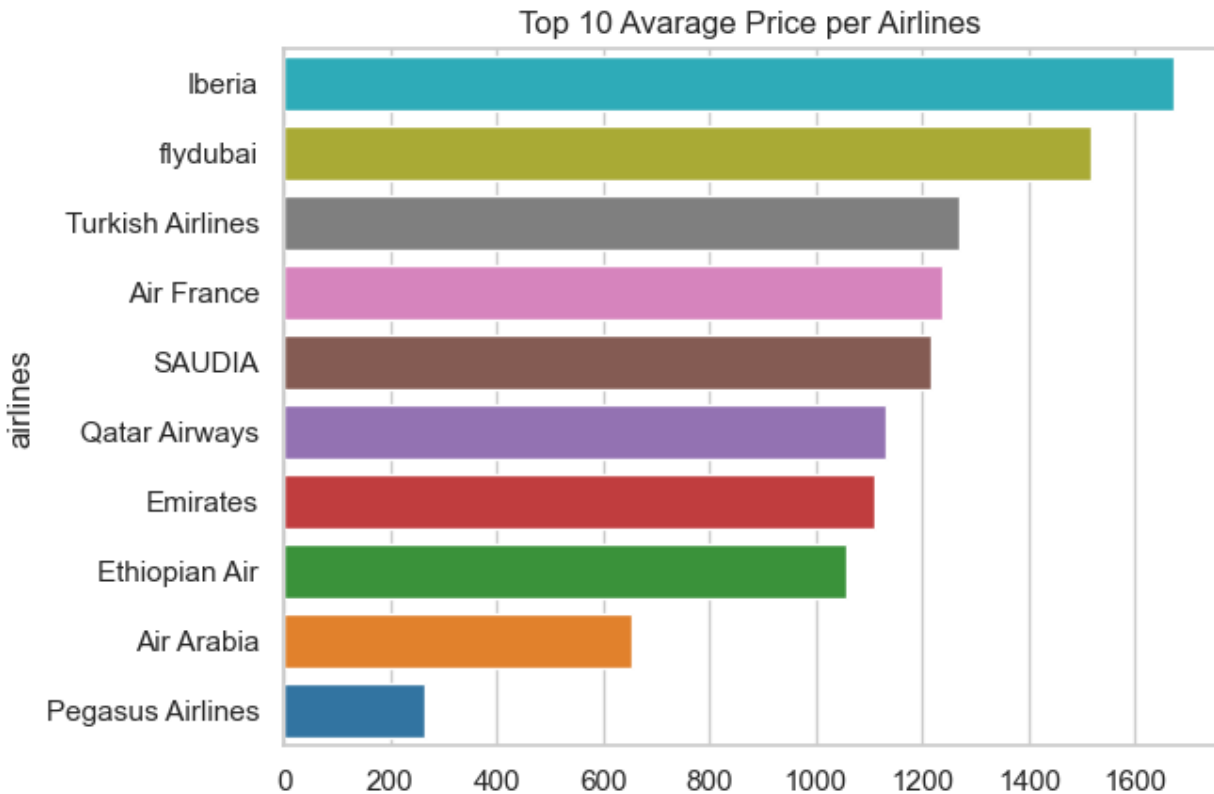
Some destinations (like SGN and TBS) have relatively lower average prices compared to others

Prices vary significantly by location, likely due to distance or airline competition.

```
top_destinations = df["airlines"].value_counts().head(10).index
top_avg_price_per_dest = df[df["airlines"].isin(top_destinations)].groupby("airlines")["price"].mean().sort_values(ascending = False)

sns.barplot(x = top_avg_price_per_dest.values, hue = top_avg_price_per_dest.values,
            y = top_avg_price_per_dest.index, palette= "tab10", legend = False)

plt.title("Top 10 Average Price per Airlines")
plt.show()
```



Airlines like Iberia and flydubai show higher average ticket prices.

Budget carriers like Wizz Air and Air Arabia offer lower average prices.

```

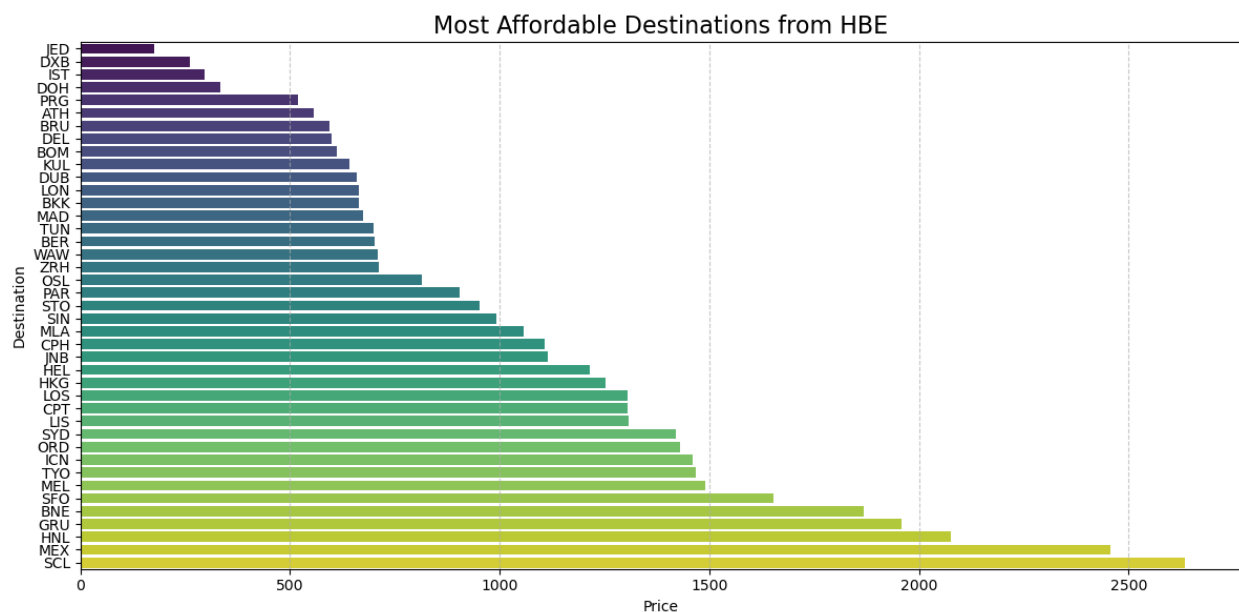
cheapest_destinations = df.groupby(['origin', 'destination'])['price'].min().reset_index()

origin_city = 'HBE'
cheapest_from_origin = cheapest_destinations[cheapest_destinations['origin'] == origin_city]

cheapest_from_origin = cheapest_from_origin.sort_values('price', ascending=True)

plt.figure(figsize=(12, 6))
sns.barplot(
    x='price',
    y='destination',
    data=cheapest_from_origin,
    palette='viridis'
)
plt.title(f'Most Affordable Destinations from {origin_city}', fontsize=16)
plt.xlabel('Price')
plt.ylabel('Destination')
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()


```



The visualization above shows — destinations like Jeddah, Dubai, and Istanbul are the most affordable from HBE(Borg El Arab)

```
correlation_matrix = df[["price", "stops", "duration"]].corr()

sns.heatmap(correlation_matrix, annot = True, cmap="coolwarm")
plt.title("Correlation Matrix")
plt.show()
```

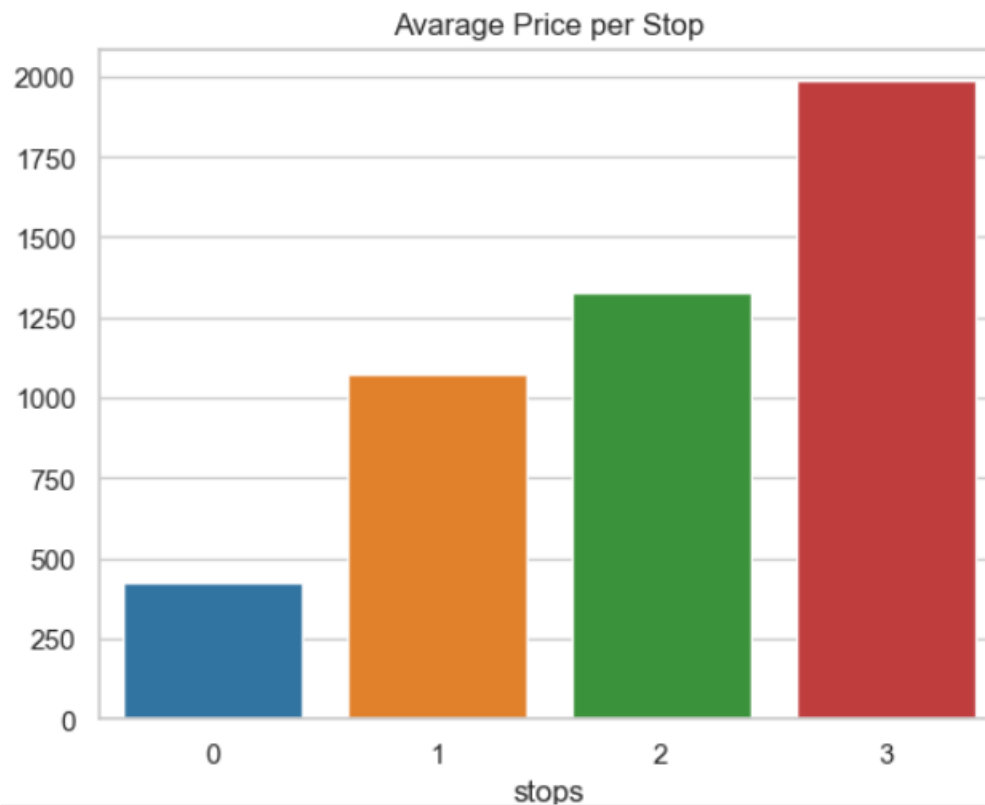


The heatmap reveals a strong correlation between duration and the number of stops, as well as a significant correlation between price and duration.

```
avg_price_per_stop = df[["stops", "price"]].groupby("stops")['price'].mean()

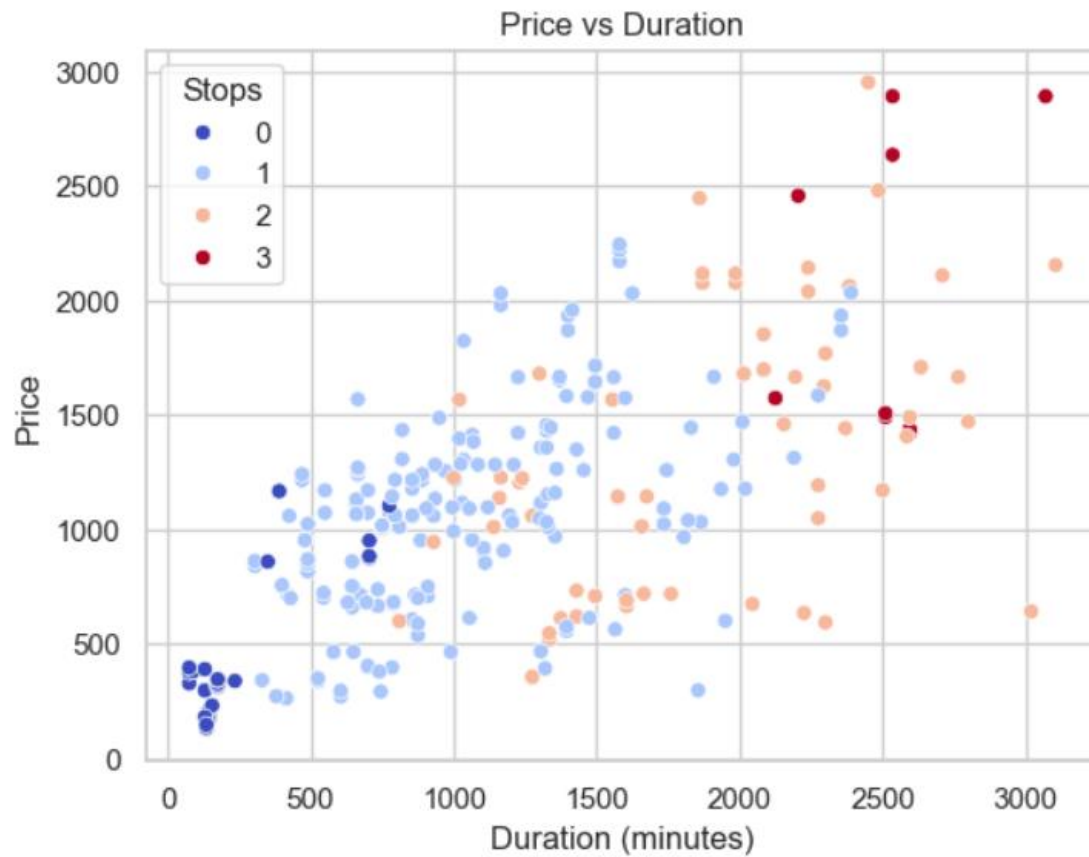
sns.barplot(x = avg_price_per_stop.index ,
            hue = avg_price_per_stop.index ,
            y = avg_price_per_stop.values ,
            palette= "tab10", legend = False)

plt.title("Avarage Price per Stop")
plt.show()
```



increasing the number of stops will increase the price

```
sns.scatterplot(data=df, x='duration', y='price', hue='stops', palette='coolwarm')
plt.title('Price vs Duration')
plt.xlabel('Duration (minutes)')
plt.ylabel('Price')
plt.legend(title='Stops')
plt.show()
```



How do prices vary by day of the week or booking time?

Could not answer because the data all in the same week

Are there noticeable seasonal trends?

Could not answer because there is not historical data

- Build interactive dashboards with Streamlit.



Data Storage:

```
1 from pymongo import MongoClient
2
3
4 def write_df_to_mongoDB( my_df, database_name = 'flights' , collection_name = 'tickets', mongodb_port = 27017):
5
6     client = MongoClient('localhost',int(mongodb_port))
7     db = client[database_name]
8     collection = db[collection_name]
9     df_cleaned = my_df.drop(columns=['Unnamed: 0'])
10    data_to_insert = df_cleaned.to_dict(orient='records')
11
12    client = MongoClient(f'mongodb://localhost:{mongodb_port}/')
13    db = client[database_name]
14    collection = db[collection_name]
15
16
17    insert_result = collection.insert_many(data_to_insert)
18    print(f"Number of document inserted: {len(insert_result.inserted_ids)}")
19    print('Done')
20    return
```

```
df = pd.read_csv("data.csv")
write_df_to_mongoDB(df,mongodb_port = 27017)
```

```
Number of document inserted: 280
Done
```

Type a query: { field: 'value' } or [Generate query](#)

Explain

Reset

Find

Options

ADD DATA

EXPORT DATA

UPDATE

DELETE

25

1 - 25 of 280

```
_id: ObjectId('689e1489ab56cff83cea01c7')
origin: "HBE"
destination: "MLA"
price: 1058.4
airlines: "Turkish Airlines"
stops: 1
duration: 425
```

```
_id: ObjectId('689e1489ab56cff83cea01c8')
origin: "HBE"
destination: "MLA"
price: 1058.4
airlines: "Turkish Airlines"
stops: 1
duration: 930
```

```
_id: ObjectId('689e1489ab56cff83cea01c9')
origin: "HBE"
destination: "PRG"
price: 519.48
airlines: "طيران الجزيرة"
stops: 2
duration: 1335
```

Expected Outcome:

An interactive dashboard showcasing insights on optimal travel times and destinations. The project offers value for cost-efficient travel planning and understanding pricing trends